

Forecasting: Exam assignment 2024

IÉSEG School of Management

Practical issues

1. The deadline for handing in the paper is 2024-03-24 (midnight). Papers that are submitted after the deadline will not be accepted. In case of discussion, the time stamp in the email will be used as a criterion (unless grave circumstances can be argued).
2. Send the paper, the R scripts, and for the second exercise also the data, by email to f.vandenbossche@ieseg.fr and p.borchert@ieseg.fr.
3. Please put your name on the report, and number the pages.

Data sets

The data for the first exercise are in the CSV file `EU_trip_expenditures.csv`, to be found on IÉSEG Online. The source and description of the dataset are available here: [Data](#), [Documentation](#).

Once you've assigned a working directory, you can load the data. For example, you can use the following code:

```
setwd("C:/Temp") # Specify your own working directory here.
data <- read_csv("EU_trip_expenditures.csv")
```

Exercise 1

The data set **Trip Expenditures** contains the average expenditure per night for EU citizens travelling within EU countries, from January 2012 to December 2022.

The dataset contains multiple observations per month for domestic and foreign trips, categorized by the length of the stay. Aggregate the data to include one observation per month representing the average trip expenditure per night across all categories.

Split the time series in a **training set up to March 2017** and a **test set from April 2017 up to March 2020**. Use the training set for estimation of the methods/models, and use the test set for assessing their forecast accuracy.

The remaining observations (April 2020 - December 2022) do not belong to the training/test set, but we keep them for later reference.

1. Explore the data using relevant graphs, and discuss the properties of the data. Include and discuss a time series plot, a seasonal plot, a seasonal subseries plot and a (P)ACF plot.
2. Discuss whether a transformation and/or any other adjustment of the time series would be useful. If so, apply the most appropriate transformation and/or adjustments. Also, report the optimal Box-Cox lambda value that could be used to transform the time series. Clarify how you will proceed with the transformation in the remainder of the exercise.
3. Create forecasts using the seasonal naive method. Check the residual diagnostics (including the Ljung-Box test) and the forecast accuracy (on the test set).
4. Use an STL decomposition to forecast the time series. Use the various underlying forecasting methods for the seasonally adjusted data (naive, rwdrift, ets, arima). Check the residual diagnostics and the forecast accuracy and select the best performing STL decomposition.
5. Generate forecasts using ETS. First select two appropriate models yourself and discuss their performance. Compare these models with the results of the automated ETS procedure. Check the residual diagnostics and the forecast accuracy for the various ETS models you've considered. Present and discuss the parameters of the final ETS model. Write the state space model representation for the final model, including the parameter estimates. Show the forecasts in a graph.
6. Generate forecasts using the `auto.arima` procedure. Write the estimated model using the backward shift operator and including the parameter estimates. Check the residual diagnostics and the forecast accuracy. Discuss your results and compare these with at least one carefully selected other candidate ARIMA model, based on a documented analysis of the ACF and PACF (e.g. check if small changes in the model specification improve the properties of the residuals and/or the forecast accuracy). Check if this candidate model results in better residual diagnostics and/or forecast accuracy.
7. Compare the different models (naive, STL, ETS, ARIMA) in terms of residual diagnostics and forecast accuracy. Present the results in a summary table. Analyse your results and select your final model.
8. Generate out of sample forecasts up to December 2023, based on the complete time series (**January 2012 - March 2020**). Present your results.
9. Now consider the last observations in the time series (April 2020 - December 2022). They include the COVID-19 pandemic. Based on the data and your forecast, what can you conclude about the impact of the pandemic on travel expenditures in the EU?
10. *Extra:* Apply TBATS and/or NNAR and compare the performance to your selected model. Explain the key differences between your selected model and TBATS or NNAR.

Exercise 2

For this exercise, find a recent and relevant time series to forecast. The time series should be original (not from R packages!), recent, sufficiently long and must include a seasonal component. Please send your selected time series with a link to the data source to f.vandenbossche@ieseg.fr and p.borchert@ieseg.fr for validation.

The time series is analysed according to a carefully selected forecasting process, using at least two techniques (other than the naive methods) that have been discussed during the lectures, and results should be compared. **Just providing ets and auto.arima results is not sufficient.** Describe your approach, and motivate your choices. Present a full analysis and description of the time series forecasting process.

General guidelines

- **Note that a fully elaborated paper is expected. The provided R-code will be checked only in case of doubt. What counts is your description, analysis and interpretation in the paper. A collection of R code and output is not considered a fully elaborated paper.**
- In each step of the exercises, discuss your results and explain your choices. Use additional tables and graphs to clarify your answer. Just providing R code and corresponding output is not sufficient to pass. You'll need to show some understanding of the course content.
- Please make sure that you analyse the correct data. This is very easy to verify in the first step of your analysis. I do not grade exercises on wrong data.
- Keep in mind that impudently copying text from the course material (including slides, R forecasting labs, solutions...) or from external sources is considered a very bad practice (it is plagiarism). Use references to sources whenever necessary. Cases of plagiarism will be reported to the program director.
- Please, read and re-read your own work before handing it in! Inconsistencies and errors make a paper very difficult to read.
- It is very useful to see the estimated parameters of your models, and your corresponding interpretation in the output. Therefore, include these in your answers when this is asked.
- Do not include useless models. If the time series is trending, then a naive (non-seasonal) forecasts or a SES are not useful, so do not include them.
- If you split the data set in a training and test set, make sure that you assess forecast accuracy on the test set, not (only) on the training set or on the complete data set.

Good Luck, and Happy Forecasting!